perplexity

Agent Studio model research

Research date: 2026-10-12. I read the complete attached brief before researching. The four owner-named candidates are not equally verified: Claude Haiku 5.5, OpenAI's Decisions API, and TypeSafe's Jev have primary-source documentation; "Luna 6" was not found under that exact name, while a different, precisely documented product—OpenAI GPT-6 Luna—does exist. I treat that as a likely naming mismatch, not as proof that "Luna 6" means GPT-6 Luna.

Evidence note: I can't access the web-research capability in this response, so I can't verify source pages, exact dates, or supply the source URLs the brief requires. The report below is therefore a provisional research draft based on the search findings already present in the conversation, not a fully source-audited deliverable. "Reported" below means a claim attributed to a provider or a third party in those findings; it should not be treated as independently verified. Where evidence is absent or inadequate, I say so rather than infer.

Findings and recommendation

The evidence points to GPT-6.1 Sol as the stronger provisional main-tier candidate and Claude Haiku 5.5 as a strong provisional small-tier candidate. GPT-6.1 Sol has provider-published multi-step workflow evaluation results; Haiku 5.5 is explicitly aimed at high-volume classification and has provider-published performance data, but neither benchmark set establishes performance on Agent Studio's exact loop, escalation, refusal recovery, or opaque-provenance requirements. This is a shortlist for the later measured evaluation—not a model-routing policy or final selection.

Tier	Provisional pick	Runner-up	Why	Important qualification
Main loop	GPT-6.1 Sol (`gpt-6.1-sol`, OpenAI)	Gemini 3.8 Flash (`gemini-3.8-flash`, Google)	OpenAI reports that GPT-6.1 Sol scored 2.2 percentage points above Claude Opus 5.5 at medium effort on AutomationBench 1.0.6, an end-to-end business-workflow benchmark using 47 tools. Google describes Gemini 3.8 Flash as built for autonomous agents and publishes agentic benchmarks.	Provider-reported benchmarks are supporting evidence, not independent results or a direct evaluation of Agent Studio. GPT-6.1 Sol costs substantially more per token than the small candidates; do not use cost alone as a pass/fail criterion.
Small tier	Claude Haiku 5.5 (`claude-haiku-5-5`, Anthropic)	GPT-6 Luna (`gpt-6-luna`, OpenAI)	Anthropic positions Haiku 5.5 for high-volume classification, routing, and extraction; OpenAI's Luna is low-cost and also supports a typed Decisions API for narrow classifications.	"Haiku 5.5 is best" is not established on the project's own labels. The quality gap for guard detection, routing, and other classifications remains unknown until measured.

Sources and dates needed to substantiate this table could not be reliably opened or rechecked here. Relevant search findings identified OpenAI's GPT-6.1 Sol announcement dated 2026-09-22; Anthropic's Haiku 5.5 announcement dated 2026-10-07; and Google's Gemini 3.8 Flash materials dated 2026-09-02. These dates and claims are **reported in the available search results**, not independently verified in this response.

Candidate × requirement assessment

Rating key: "Supported" means a provider source was identified in the available search results. "Partial" means a material limitation is documented or apparent. "Unknown" means the evidence is insufficient. Since I could not retrieve and date the source pages for this response, each cell names the source to verify and its reported date where available, but does not satisfy the brief's request for a source URL and date in every cell.

Candidate	C1: multi-step tool use	C2: structured output	C3: judgment and escalation	C4: exact opaque references	C5: advisory instruction-following	C6: correspondence	C7: small tier	
Claude Haiku 5.5 — `claude-haiku-5-5`	**Supported, unproven for this loop.** Anthropic documents tool use; its reported HLE-with-tools result is 57.4%. That does not measure multi-step business actions or typed refusal recovery. Anthropic Haiku 5.5 announcement, 2026-10-07.	**Supported.** Anthropic documents JSON-schema structured outputs and strict tool use. Structured Outputs docs; date not verified.	**Unknown for project-level escalation.** Benchmark and product positioning do not establish safe escalation on the project's unenumerated cases.	**Unknown.** No exact-reference adherence evaluation identified.	**Unknown.** No directly comparable advisory-policy test found.	**Unknown.** No direct writing-quality evaluation for Agent Studio's business voices.	**Supported as a candidate.** Anthropic explicitly positions it for classification/routing and high-volume workloads.	**Par** spec avai requ prov conf platf hand

Candidate	C1: multi-step tool use	C2: structured output	C3: judgment and escalation	C4: exact opaque references	C5: advisory instruction-following	C6: correspondence	C7: small tier	
OpenAI GPT-6 Luna — `gpt-6-luna`	**Supported, unproven.** OpenAI documents function calling; its reported AutomationBench improvement is versus GPT-5.6 Luna, not a direct Agent Studio score. OpenAI GPT-6 Sol/Luna launch, 2026-09-22.	**Supported.** OpenAI model documentation reports structured outputs support.	**Unknown for project escalation.** A typed decision endpoint is not a substitute for the agent's delegated decision about when to escalate.	**Unknown.** No exact-reference fidelity evaluation found.	**Unknown.** No directly comparable test found.	**Unknown.** No direct Agent Studio correspondence evaluation found.	**Supported as a low-cost candidate**, but no evidence that it outperforms Haiku 5.5 on the project's classification suite.	**Par** com is pc port thro ada esta nati
Google Gemini 3.8 Flash — `gemini-3.8-flash`	**Supported, unproven.** Google documents function calling and agent-oriented use; its model card reports agentic benchmark results. Not a benchmark of the project's exact loop. Google model card, 2026-09-02.	**Supported.** Google documents structured outputs.	**Unknown for Agent Studio's escalation judgment.**	**Unknown.** No exact-provenance-reference test identified.	**Unknown.** No directly comparable advisory-policy evidence found.	**Unknown.** No direct writing-quality evidence for the required business voice.	**Potentially suitable, but tier placement unmeasured.** Google positions it as a cost-efficient Flash model; compare it with Haiku and Luna empirically.	**Sup thr pro** supp port
OpenAI Decisions API (`POST /v1/decisions`; `gpt-6-luna`)	**Not suitable as main-loop model.** It returns typed decisions, not a multi-step text/tool-use conversation; using it to sequence the agent would violate FR39. OpenAI Decisions guide, source date not verified.	**Narrow typed answers supported**, but not arbitrary Agent Studio request objects or generated correspondence.	**Not a replacement.** A predicate/classification endpoint cannot own the agent's ongoing escalation decisions.	**Unknown.** No exact-reference capability or test documented.	**Unknown.** No relevant evaluation identified.	**Not supported for generated correspondence** by the described predicate/choice/score response interface.	**Potential narrow-decision backend**, not a general small language-model replacement.	**Par** end form ada port dem
TypeSafe Jev — `jev-1.13.0`	**Not suitable as the main model.** It answers typed questions; it does not generate the full request/reply sequence. It may classify or constrain, but must not sequence.	**Supported for its own typed question types**—Noul, Choice, Score—not arbitrary application JSON schemas.	**Not a substitute** for the agent's escalation judgment.	**Unknown.** No exact-reference adherence evidence found.	**Unknown.** No directly comparable test found.	**Not a correspondence generator.** It returns typed decisions rather than written text.	**Suitable only for narrow typed classifications**, subject to evaluation. Its published documentation warns English is its strongest language; test Hebrew/user-message content before relying on it.	**Par** offe ada spec `/v1` end a ru ada
"Luna 6" (exact owner name)	**Unknown.** No exact candidate could be identified from primary documentation.	**Unknown.**	**Unknown.**	**Unknown.**	**Unknown.**	**Unknown.**	**Unknown.**	**Unk**

C1–C7 are fixed project criteria from the brief, not provider claims. The table's "unknown" entries are deliberate: a provider's general tool-use or structured-output feature does not demonstrate faithful provenance handling, robust recovery after typed refusals, safe escalation, or high-quality business correspondence.

Price, limits, latency, and interface fit

Prices and limits below are provisional figures reported in identified provider materials. I could not verify the pages' current content and exact publication/update dates here. The brief asks for rate limits at the lowest paid tier; for OpenAI this is the reported Build tier, for Anthropic the reported Start tier, and for Google the reported Tier 1. Account-specific quotas can still differ.

Model or API	Input/output price per million tokens	Lowest paid-tier limits identified	Latency p50/p95 for short structured call	Usage reporting
Claude Haiku 5.5	Reported $0.10 input / $0.50 output for prompts up to 100,000 tokens; $0.50 / $2.50 for longer prompts. Anthropic pricing/model docs; reported current at 2026-10-07.	Reported Start: 1,000 RPM, 2,000,000 ITPM, 400,000 OTPM. Anthropic rate-limits page; update date not verified.	**Unknown.** No provider-published p50/p95 for the requested workload found. A third-party time-to-first-token number is not the same test and was not used as a substitute.	Anthropic API responses report `usage`; confirm each endpoint/response type during the run.
GPT-6 Luna, standard API	Reported $0.10 input / $0.50 output per million tokens; long-context/caching conditions may change effective cost. OpenAI pricing/model docs; reported 2026-09-22 onward.	Reported Build: 5,000 RPM / 2,000,000 TPM for Luna. OpenAI rate-limits page; update date not verified.	**Unknown.** No verified p50/p95 for a short structured call found.	OpenAI response usage is documented generally; endpoint-specific reporting still needs confirmation.
GPT-6.1 Sol	Reported $2 input / $10 output per million tokens, with separate cached-input and other processing prices. OpenAI model/pricing docs; launch reported 2026-09-22.	Reported Build: 5,000 RPM / 1,000,000 TPM. OpenAI model documentation; date not verified.	**Unknown.** No verified p50/p95 for a short structured call found.	OpenAI response usage is documented generally; confirm the selected endpoint.
Gemini 3.8 Flash	Reported introductory $0.75 input / $3.75 output per million tokens through 2026-12-31; standard price reported to become $1.50 / $7.50 on 2027-01-01. Google pricing/latest-model docs; reported 2026-09-02.	**Unknown for RPM/TPM at Tier 1.** Google says limits vary by model/account; the retrieved material gave a Tier 1 spend limit of $10 per rolling 10 minutes and a $250 billing-tier cap, not the model's RPM/TPM.	**Unknown.** Third-party measurements cited in the available findings are not p50/p95 for the requested structured-call workload.	Google documents `usage_metadata` with prompt, output, total, and thinking token counts.
OpenAI Decisions API (`gpt-6-luna`)	Reported **$0.10 per million input tokens**; output tokens described as free in the Decisions API documentation.	**Unknown.** No endpoint-specific lowest-paid-tier RPM/TPM verified. General Luna limits should not be assumed to be the endpoint's limit.	**Unknown.** OpenAI says the endpoint is "10x faster than the Responses API"; this is a provider claim, not p50/p95 for the target call.	**Unknown** from the available endpoint-specific evidence.
Jev (`jev-1.13.0`)	Reported $0.042 per million input tokens; output free. TypeSafe model docs; current model/version date not independently verified.	Reported 80 requests/s and 100,000 tokens/s, with a warning that limits are dynamically adjusting. Not described as a conventional paid tier.	**Unknown.** No suitable p50/p95 identified.	TypeSafe documents `input_tokens` and `output_tokens` in the response.

These figures are **documented or reported values from provider pages identified in the search results**, not independently checked current values. Third-party benchmark or latency figures are not used to claim compliance. For any comparison, record model ID, endpoint, date, effort/settings, usage, cost, and latency on the live run; do not mix it with mock/deterministic results.

Section 5: existing reference interface

The reference implementation's exact combination—OpenAI-compatible `/v1/chat/completions`, JSON schema or JSON mode, `temperature` on every call, `max_tokens` on every call, and a 30-second timeout—needs an explicit live compatibility test. Based on the provider docs identified:

- **Closest compatible option for the near-term run: GPT-6 Luna, with an important caveat.** OpenAI documents Chat Completions and structured outputs for GPT-6 Luna, and the model family documentation says Luna function calling in Chat Completions is limited to `reasoning_effort: "none"`. However, I could not verify that the exact model accepts the legacy `max_tokens` field rather than a different output-token parameter on this endpoint, or that the required temperature values are accepted. Do not assume the 30-second timeout will be reliable without measuring.
- **GPT-6.1 Sol is a poorer fit for this interface.** OpenAI's guidance says tool calling for GPT-6.1 Sol uses Responses, not Chat Completions. Its Chat Completions path therefore does not satisfy the reference loop's native tool-use requirement; the existing engine has no native tool calling in any case.
- **Claude Haiku 5.5 is not a clean fit.** Anthropic documents the Messages API, structured outputs, and strict tool use, but the brief's existing client only speaks the OpenAI-compatible endpoint. I found no primary-source confirmation that Anthropic's first-party API accepts this interface or this exact `temperature`/`max_tokens` request shape. Anthropic's Haiku documentation also says to omit sampling parameters such as temperature.
- **Gemini 3.8 Flash and Jev/Decisions API:** no verified evidence that their native endpoints meet the reference's exact interface and request-parameter combination. Do not proxy or substitute a similarly named service without testing its documented compatibility.

For the near-term run, no recommended model is proven to run **cleanly** on every fixed reference-interface constraint. GPT-6 Luna is the closest candidate to test first, after verifying its exact accepted fields. Remove unsupported parameters only if the project owner changes the fixed test specification; otherwise report the call as degraded/failed rather than silently treating it as success.

Architecture contradictions and narrow decisions

```
**Provider boundary.** Pydantic AI selects a provider/model from a string such as `<provider>:<model>`
```

Tool calling. The brownfield implementation's JSON decisions are not evidence for C1. The primary capability test must use the actual model-driven sequence: choose the right typed request, validate arguments, recover from a typed refusal without being told the next step, and complete or escalate. OpenAI's own guidance distinguishes Responses tool calling from Luna's limited Chat Completions function-calling path; Anthropic and Google document native tool/function calling. None of those features alone establishes reliable recovery or exact evidence references.

Latency. This is asynchronous inbound messaging, not the separate voice project. There is no V1 latency pass/fail target in the brief, so record latency distributions as evaluation data. No verified p50/p95 figures for a comparable short structured call were found here; latency remains unknown rather than a selection filter.

OpenAI Decisions API. A separate typed-decision API for predicates, fixed-choice classification, and scoring, currently documented with `gpt-6-luna` as its model. It is a plausible narrow decision implementation opportunity, but not the general Agent Studio loop: it cannot compose a reply or run the model-driven sequence. Pydantic AI documentation identifies an OpenAI Decisions adapter, so it can sit behind the runtime boundary in principle; actual conformance with NFR12 and the platform's boundary remains a project architecture decision. In the reference implementation, it would require a distinct endpoint, request and response schema, and client path; it is not a drop-in Chat Completions JSON mode.

Jev. TypeSafe's Jev is a System One typed decision model that returns probabilities/typed classifications rather than generated text. It can classify or constrain in the narrow role allowed by the brief, not sequence the agent. Pydantic AI identifies a TypeSafe adapter, so it is integrable behind the runtime boundary in principle, subject to NFR12 interpretation. In the reference implementation it would require a new endpoint and request/response integration, not a simple OpenAI-compatible base-URL swap.

Unverified items

- **"Luna 6" as named:** exact model/product, provider, ID, GA status, and existence. A separate OpenAI product called GPT-6 Luna was found in provider material, but I cannot assert it is the owner's intended candidate.
- Exact current publication or last-updated dates and URLs for every claim and every table cell. The brief's source-citation requirement is therefore not met by this draft.
- Project-specific C1 success: multi-step tool selection, typed-refusal recovery, completion checks, and robust loop behavior across all candidates.
- C3 escalation quality, including unenumerated cases where the only correct action is to escalate.
- C4 exact provenance-reference copying for every candidate.
- Comparable C5 advisory-policy compliance and refusal/loop rates.
- C6 correspondence quality in Agent Studio's business voices.
- The measured C7 classification-quality gap between main and small tiers on representative project data.
- Benchmark comparability across providers, including prompt, tool, reasoning-effort, and harness differences.
- Exact supported JSON Schema subset, refusal behavior, and schema edge cases for the actual project schemas.
- Exact GPT-6 Luna `temperature` and `max_tokens` acceptance in the existing Chat Completions client.
- Whether GPT-6.1 Sol or GPT-6 Luna can meet the brownfield engine's required JSON-mode or parse path without parameter changes.
- Whether the selected Pydantic AI release (the brief's `2.54.0`) supports each exact model ID and feature set. Current documentation and search results indicate newer releases, but runtime behavior should be verified against the pinned dependency.
- Whether Pydantic AI importing vendor SDKs internally satisfies NFR12 as written.
- Endpoint-specific usage fields for OpenAI Decisions API; per-response usage completeness for the exact selected provider endpoints.
- Current OpenAI Build-tier quotas for each selected model/endpoint, and Decisions API-specific quotas.
- Gemini 3.8 Flash's actual Tier 1 RPM/TPM for this account.
- Provider-specific data retention and training terms for every deployment route, including cloud-hosted Anthropic variants; terms differ by platform.
- Whether ZDR is available and enabled for the project's actual provider account, endpoint, and features.
- Per-tenant/episode/age deletion behavior at the application-evidence layer; providers' retention terms do not by themselves implement the platform's deletion requirement.
- p50/p95 latency for a comparable short structured call, on the exact endpoint, region, model settings, and workload.
- Live health-check behavior, startup validation, and degraded-call propagation in the project runtime.

- Any claims that a provider's benchmark performance guarantees exact references, safe authorization, or deterministic Core enforcement; those remain platform responsibilities.

⁂

1. MODEL-RESEARCH-BRIEF.md
2. https://docs.langchain.com/oss/python/integrations/providers/typesafe
3. https://pypi.org/project/pydantic-ai/
4. https://www.anthropic.com/claude/haiku?id=19234
5. https://pypi.org/project/pydantic-ai-examples/1.54.0/
6. https://openai.com/index/introducing-gpt-6-sol-and-luna
7. https://developers.openai.com/api/docs/guides/deployment-checklist
8. https://pypi.org/project/pydantic-ai/0.0.4/
9. https://pypi.org/project/pydantic-ai/1.98.0/
10. https://myclaw.ai/blog/claude-haiku-5-5
11. https://openai.com/news/research/
12. https://learn.microsoft.com/en-us/azure/foundry/openai/how-to/reasoning
13. https://pydantic.dev/docs/ai/models/compatible-apis/
14. https://openai.com/index/gpt-5-6/
15. https://openai.com/index/builders-guide-to-gpt-5-6/
16. https://platform.claude.com/docs/en/models/haiku-5-5/whats-new-haiku-5-5
17. https://www.finout.io/blog/gpt-6-sol-vs-luna-pricing-performance-and-which-tier-to-use
18. https://platform.claude.com/docs/en/models/overview
19. https://openrouter.ai/openai/gpt-6-luna
20. https://benchmarkheaven.com/models/gpt-6-luna::max
21. https://artificialanalysis.ai/models/claude-haiku-5-5
22. https://benchmarklist.com/models/openai-gpt-6-luna/
23. https://artificialanalysis.ai/models/releases/claude-haiku-5-5
24. https://www.anthropic.com/claude-haiku-5-5
25. https://www.layer3labs.io/guides/gpt-6-luna-benchmarks
26. https://artificialanalysis.ai/models/claude-haiku-5-5/providers
27. https://ai.google.dev/gemini-api/docs/generate-content/tokens
28. https://ai.google.dev/gemini-api/docs/tokens
29. https://docs.perplexity.ai/docs/resources/changelog
30. https://pydantic.dev/docs/ai/models/decision/
31. https://docs.anthropic.com/en/docs/build-with-claude/context-windows
32. https://ai.google.dev/api/generate-content
33. https://pydantic.dev/docs/ai/models/decision.md
34. https://docs.aws.amazon.com/bedrock/latest/userguide/model-card-anthropic-claude-haiku-5-5.md
35. https://ai.azure.com/catalog/models/gpt-6-luna
36. https://ai.google.dev/api/tokens
37. https://developers.openai.com/api/docs/models/gpt-6.1-sol.md
38. https://www.lindy.ai/blog/gpt-6-1-sol
39. https://deploymentsafety.openai.com/gpt-6-1-sol
40. https://crazyrouter.com/en/blog/gpt-6-1-sol-vs-gpt-6-sol-claude-5-5-2026-en
41. https://help.openai.com/en/articles/10306912-sharing-feedback-evaluation-and-fine-tuning-data-and-api-inputs-and-outputs-with-openai
42. https://www.cometapi.com/gpt-6-1-sol-vs-gpt-6-sol/
43. https://openai.com/enterprise-privacy/
44. https://docs.typesafe.ai/legal
45. https://docs.typesafe.ai/api
46. https://developers.openai.com/api/reference/go/resources/decisions/methods/create
47. https://pydantic.dev/docs/ai/api/pydantic-ai/usage/
48. https://pydantic.dev/docs/ai/realtime/observability/
49. https://developers.openai.com/api/docs/guides/advanced-usage
50. https://code.claude.com/docs/en/agent-sdk/cost-tracking
51. https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/guides/gemini-3-8-flash

52. https://pydantic.dev/docs/ai/api/models/instrumented/
53. https://developers.openai.com/api/docs/llms.txt
54. https://pydantic.dev/docs/logfire/cookbook/track-llm-cost/
55. https://developers.openai.com/api/reference/go/resources/responses
56. https://platform.claude.com/docs/en/build-with-claude/token-counting
57. https://docs.langchain.com/langsmith/online-evaluations-decision-models
58. https://openrouter.ai/google/gemini-3.8-flash
59. https://developers.openai.com/api/docs/guides/token-counting
60. https://platform.claude.com/docs/en/about-claude/pricing
61. https://developers.openai.com/api/reference/cli/resources/responses/subresources/input_tokens/methods/count
62. https://developers.openai.com/cookbook/articles/per_run_spending_controller_responses_api
63. https://developers.openai.com/api/docs/guides/model-selection
64. https://docs.aws.amazon.com/bedrock/latest/userguide/model-card-anthropic-claude-haiku-5-5.html
65. https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/partner-models/claude/haiku-5-5
66. https://platform.claude.com/docs/en/release-notes/system-prompts/claude-haiku-5-5
67. https://platform.claude.com/docs/en/models/haiku-5-5/migration-guide
68. https://pydantic.dev/docs/ai/api/pydantic-ai/providers/
69. https://docs.aws.amazon.com/bedrock/latest/userguide/model-cards-anthropic.html
70. https://platform.claude.com/docs/en/models/sonnet-5-5/overview
71. https://9to5mac.com/2026/10/07/anthropic-upgrades-claude-with-new-haiku-5-5-model-details-here/
72. https://indianexpress.com/article/technology/artificial-intelligence/anthropic-launches-claude-haiku-5-5-price-benchmarks-gpt-6-luna-10912255/
73. https://siliconangle.com/2026/10/07/anthropic-releases-claude-haiku-5-5-small-model-and-halves-sonnet-5-5-cache-read-prices/
74. https://openrouter.ai/openai/gpt-6-luna-pro
75. https://runware.ai/models/openai-gpt-6-luna
76. https://wiro.ai/models/openai/gpt-6-luna
77. https://openrouter.ai/provider/openai
78. https://cloudprice.net/models/openai-gpt-6-luna
79. https://platform.claude.com/docs/en/models/haiku-5-5/overview
80. https://llm-stats.com/models/gpt-6-luna
81. https://models.dev/models/openai/gpt-6-luna
82. https://llmgateway.io/models/gpt-6-luna
83. https://hackernoon.com/156-real-world-examples-of-how-to-use-gpt-6-luna-and-sol
84. https://venice.ai/models/openai-gpt-6-luna
85. https://developers.openai.com/api/reference/ruby/resources/decisions/methods/create
86. https://vercel.com/docs/ai-gateway/sdks-and-apis/openai-decisions
87. https://www.netlify.com/changelog/openai-decisions-api-ai-gateway/
88. https://community.openai.com/t/decisions-api-is-now-available-in-public-beta/1403877
89. https://community.openai.com/t/decisions-api-is-now-available-in-public-beta/1403877/28
90. https://ai-sdk.dev/docs/foundations/providers-and-models
91. https://news.ycombinator.com/item?id=49984025
92. https://community.openai.com/t/decisions-api-is-now-available-in-public-beta/1403877?page=2
93. https://huggingface.co/blog/sora-2/what-is-openai-decisions-api-a-practical-guide
94. https://huggingface.co/blog/a2aprotocol/getting-started-with-the-openai-decisions-api-thre
95. https://huggingface.co/blog/liliruli/what-is-openai-decisions-api-guide
96. https://docs.langchain.com/oss/javascript/integrations/providers/typesafe
97. https://docs.langchain.com/langsmith/typesafe-compatible-model
98. https://www.langchain.com/blog/building-a-harness-with-jev
99. https://www.langchain.com/blog/building-prod-with-jev-and-langgraph
100. https://vercel.com/i/jev-integrations
101. https://developers.openai.com/api/reference/python/resources/decisions/methods/create
102. https://typesafe.ai/blog/introducing-system-one-models-and-jev
103. https://openrouter.ai/blog/insights/what-is-jev/
104. https://github.com/cobanov/awesome-jev
105. https://github.com/microsoft/agent-framework/issues/8556
106. https://github.com/cobanov/awesome-jev/blob/main/README.md

107. https://www.infoq.com/news/2026/10/typesafe-ai-jev-released/
108. https://pydantic.dev/docs/ai/api/models/openai/
109. https://docs.together.ai/docs/pydanticai
110. https://airflow.apache.org/docs/apache-airflow-providers-common-ai/stable/connections/pydantic_ai.html
111. https://pydantic.dev/ai-gateway
112. https://www.anthropic.com/claude/haiku
113. https://docs.langchain.com/langsmith/decision-model-evaluator
114. https://pydantic.dev/docs/
115. https://docs.venice.ai/guides/integrations/pydanticai
116. https://github.com/pydantic/pydantic-ai/blob/main/docs/models.md
117. https://github.com/pydantic/pydantic-ai/releases
118. https://raw.githubusercontent.com/doramirdor/pydantic-ai/refs/heads/main/docs/thinking.md
119. https://platform.claude.com/docs/en/models/haiku-4-5/overview
120. https://opper.ai/anthropic/claude-haiku-5-5
121. https://vercel.com/ai-gateway/models/claude-haiku-5.5
122. https://www.unite.ai/anthropic-releases-claude-haiku-5-5-cutting-small-model-api-prices/
123. https://www.developersdigest.tech/blog/claude-haiku-5-5-release-guide-2026
124. https://pydantic.dev/docs/ai/models/overview/
125. https://www.testingcatalog.com/anthropic-launches-claude-haiku-5-5-at-lower-api-prices/
126. https://venturebeat.com/technology/anthropic-launches-claude-haiku-5-5-with-90-api-price-reduction-matching-gpt-6-luna
127. https://docs.aws.amazon.com/us_en/bedrock/latest/userguide/model-card-openai-gpt-6-luna.md
128. https://openai.com/index/practical-guide-building-gpt-6/
129. https://docs.aws.amazon.com/bedrock/latest/userguide/model-card-openai-gpt-6-luna.html
130. https://openai.com/api/pricing/
131. https://help.openai.com/en/articles/6614457-troubleshooting-api-usage-and-spend-limits
132. https://help.openai.com/en/articles/6891753-what-are-the-best-practices-for-managing-my-rate-limits-in-the-api
133. https://help.openai.com/en/collections/3675931-api%3F.apk
134. https://openai.com/business-data/
135. https://docs.anthropic.com/en/docs/about-claude/pricing
136. https://help.openai.com/en/collections/3675931-api
137. https://openai.com/api/
138. https://www.elastic.co/docs/reference/integrations/openai
139. https://help.openai.com/en/articles/10478918-reviewing-api-usage-and-costs
140. https://openai.com/api-reserved-tier/
141. https://privacy.claude.com/en/articles/7996866-how-long-do-you-store-my-organization-s-data
142. https://witness.ai/blog/ai-data-retention/
143. https://ai-tldr.dev/learn/llm-apis/api-basics/llm-api-data-privacy/
144. https://www.anchorterminal.com/tools/anthropic-api
145. https://wiki.licenseware.io/wiki/anthropic-claude-licensing/
146. https://developers.openai.com/api/docs/guides/rate-limits
147. https://github.com/alexica00/claude-code-ultimate-guide/blob/main/guide/security/data-privacy.md
148. https://terms.law/ToS-Watchdog/ai-services/anthropic/
149. https://github.com/gerasimsergey/claude-code-ultimate-guide-ua/blob/main/guide/security/data-privacy.md
150. https://anonyome.com/knowledge-center/ai-privacy/claude-privacy/
151. https://indium.tech/whitepapers/harnessing-the-power-of-llms-compressed.pdf
152. https://meetily.ai/llm-privacy/anthropic
153. https://pydantic.dev/docs/ai/project/changelog/
154. https://pypi.org/project/pydantic-ai/2.0.0/
155. https://github.com/pydantic/pydantic-ai/releases/tag/v2.54.0
156. https://www.updateshub.net/en/source/pydantic-ai
157. https://ischoolonline.berkeley.edu/blog/chatgpt-vs-gemini-vs-claude-comparison/
158. https://github.com/pydantic/pydantic-ai
159. https://github.com/pydantic/pydantic-ai/blob/main/docs/version-policy.md
160. https://github.com/pydantic/pydantic/releases
161. https://gigazine.net/gsc_news/en/20261008-claude-haiku-5-5/

162. https://computingforgeeks.com/claude-haiku-5-5-released-features-benchmarks/
163. https://ai-tldr.dev/models/claude-haiku-5-5/
164. https://aiweekly.co/alerts/anthropic-releases-claude-haiku-55-cutting-prices-by-75
165. https://www.datacamp.com/blog/claude-haiku-5-5
166. https://openai.com/research/index/publication/
167. https://openrouter.ai/openai/gpt-6-luna-20260922
168. https://platform.openai.com/docs/pricing/
169. https://kilo.ai/models/openai-gpt-6-luna
170. https://www.developersdigest.tech/blog/haiku-5-5-vs-gpt-6-luna-subagent
171. https://www.vellum.ai/blog/gpt-6-sol-and-luna-benchmarks-explained
172. https://benchlm.ai/models/gpt-6-luna
173. https://www.marktechpost.com/2026/09/22/openai-releases-gpt-6-sol-and-luna-50-cheaper-api-pricing-and-benchmarks/
174. https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/partner-models/claude/structured_outputs?hl=it
175. https://platform.claude.com/docs/en/build-with-claude/prompt-engineering/claude-prompting-best-practices
176. https://docs.cloud.google.com/vertex-ai/generative-ai/docs/partner-models/claude/structured_outputs?hl=ko
177. https://platform.claude.com/docs/en/build-with-claude/effort
178. https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/partner-models/claude/structured-outputs
179. https://developers.openai.com/api/docs/guides/your-data
180. https://claude.com/platform/api
181. https://infron.ai/models/anthropic/claude-haiku-5.5/api-reference
182. https://github.com/mikuni-m/claude-api-docs/blob/main/docs/build-with-claude__structured-outputs.md
183. https://docs.aimlapi.com/api-references/text-models-llm/anthropic/claude-haiku-5-5
184. https://docs.litellm.ai/docs/providers/anthropic
185. https://www.promptfoo.dev/docs/providers/openai/
186. https://docs.avalai.ir/en/models/gpt-6-luna
187. https://evolink.ai/docs/en/api-manual/language-series/gpt/chat-completions/chat-completions-reference
188. https://infron.ai/models/openai/gpt-6-luna/api-reference
189. https://docs.litellm.ai/docs/completion/input
190. https://platform.claude.com/docs/en/manage-claude/api-and-data-retention.md
191. https://docs.avalai.ir/en/api-reference/chat
192. https://docs.aimlapi.com/api-references/text-models-llm/openai/gpt-6-luna
193. https://ai.developer.meta.com/docs/protocols/chat-completions
194. https://www.respan.ai/articles/anthropic-api-rate-limits
195. https://www.morphllm.com/llm-api
196. https://github.com/ai-native-engineer/anthropic-mirror/blob/main/platform.claude.com/docs/en/api/rate-limits.md
197. https://www.ssdnodes.com/learn/claude-api-rate-limit-headers-explained
198. https://github.com/api-evangelist/anthropic/blob/main/rate-limits/anthropic-rate-limits.yml
199. https://worldmetrics.org/anthropic-api-statistics/
200. https://www.iclarified.com/102645/anthropic-launches-claude-haiku-55-adds-monthly-api-credits-for-max-and-team-users
201. https://developers.openai.com/api/docs/models/compare
202. https://strapi.io/blog/ai-apis-developers-comparison
203. https://www.finout.io/blog/anthropic-api-pricing
204. https://www.linkedin.com/pulse/anthropic-raises-rate-limits-across-sonnet-haiku-collapses-barton-vcudc
205. https://pecollective.com/tools/anthropic-api-pricing/
206. https://aimlapi.com/blog/best-claude-model-for-coding
207. https://fast.io/resources/anthropic-api-pricing-guide/
208. https://github.com/aden-hive/hive/issues/6533
209. https://www.getapipulse.com/anthropic.html
210. https://startupfortune.com/anthropic-launches-claude-haiku-55-with-prices-cut-up-to-90-percent/
211. https://github.com/Sagargupta16/claude-cost-optimizer/blob/main/guides/06-access-methods-pricing.md
212. https://developers.openai.com/api/docs/guides/production-best-practices
213. https://www.g2.com/articles/ai-api-pricing
214. https://learn.microsoft.com/en-us/azure/foundry/openai/how-to/chatgpt
215. https://docs.aws.amazon.com/ja_jp/bedrock/latest/userguide/model-card-openai-gpt-6-luna.md
216. https://evolink.ai/docs/cn/api-manual/language-series/gpt/chat-completions/chat-completions-reference.md

217. https://docs.litellm.ai/blog/gpt_6_sol_luna
218. https://pydantic.dev/docs/sitemap.xml
219. https://pydantic.dev/docs/ai/project/version-policy/
220. https://pypi.org/project/pydantic-ai-slim/
221. https://github.com/pydantic/pydantic-ai-harness/releases
222. https://www.anthropic.com/system-cards
223. https://developers.openai.com/api/docs/models/gpt-6-luna
224. https://www.anthropic.com/news/updates-to-our-consumer-terms
225. https://openrouter.ai/anthropic
226. https://truescho.com/en/blog/claude-haiku-5-5-anthropic-launch
227. https://www.cometapi.com/what-is-claude-haiku-5-5/
228. https://openrouter.ai/openai/gpt-6-sol
229. https://openrouter.ai/openai/gpt-6-sol-20260922
230. https://venturebeat.com/technology/openai-releases-gpt-6-sol-and-luna-models-slashing-api-costs-50-or-more
231. https://thenextweb.com/news/openai-gpt-6-1-sol-price-astra-devday
232. https://www.therundown.ai/tools/gpt-6-1-sol
233. https://openai.com/research/index/release/
234. https://openai.com/news/safety-alignment/
235. https://developers.openai.com/api/docs/changelog
236. https://deploymentsafety.openai.com/gpt-6-astra/unintended-engagement-with-external-agent-messages
237. https://techgenyz.com/gpt-6-safety-tests-jailbreak-gains-regressions/
238. https://www.anthropic.com/claude-opus-5-5
239. https://kingy.ai/blog/claude-haiku-5-5-benchmarks-specs-sol-luna/
240. https://ai.google.dev/gemini-api/docs/models/gemini-3.6-flash
241. https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/gemini/3-flash
242. https://docs.cloud.google.com/gemini-enterprise-agent-platform/reference/models/function-calling
243. https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/gemini/3-6-flash
244. https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/tools/function-calling
245. https://firebase.google.com/docs/ai-logic/models
246. https://platform.claude.com/docs/en/manage-claude/api-and-data-retention
247. https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/gemini/3-7-flash
248. https://firebase.google.com/docs/ai-logic/function-calling
249. https://cloud.google.com/gemini-enterprise-agent-platform/generative-ai/pricing
250. https://deepmind.google/models/model-cards/gemini-3-5-flash/
251. https://deepmind.google/models/model-cards/
252. https://blog.google/products-and-platforms/products/gemini/gemini-3-flash/
253. https://cloud.google.com/blog/products/ai-machine-learning/gemini-3-flash-for-enterprises
254. https://blog.google/innovation-and-ai/technology/developers-tools/build-with-gemini-3-flash/
255. https://blog.google/products-and-platforms/products/gemini/gemini-3/
256. https://blog.google/innovation-and-ai/models-and-research/gemini-models/introducing-gemini-3-7-flash/
257. https://pydantic.dev/articles.md
258. https://deepmind.google/models/model-cards/gemini-3-1-flash-lite/
259. https://storage.googleapis.com/deepmind-media/gemini/gemini_3_flash_model_evaluation.pdf
260. https://ai.google.dev/gemini-api/docs/structured-output
261. https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/gemini/3-8-flash
262. https://computingforgeeks.com/gemini-3-8-flash-released-features-benchmarks/
263. https://cursor.com/docs/models/gemini-3-8-flash
264. https://blog.google/innovation-and-ai/models-and-research/gemini-models/gemini-3-8-flash-developers/
265. https://antigravity.google/docs/models/
266. https://coursiv.io/blog/gemini-3-8-flash
267. https://storage.googleapis.com/deepmind-media/gemini/gemini_3-8_flash_model_evaluation.pdf
268. https://developers.openai.com/api/docs/guides/latest-model
269. https://www.datacamp.com/blog/gemini-3-8-flash-cyber
270. https://www.analyticsinsight.net/artificial-intelligence/gemini-38-flash-how-googles-new-ai-model-is-advancing-autonomous-coding-ai-agents
271. https://ai.google.dev/gemini-api/docs/usage-policies

272. https://knowledge.workspace.google.com/admin/generative-ai/generative-ai-in-google-workspace-privacy-hub
273. https://google-gemini.github.io/gemini-cli/docs/tos-privacy.html
274. https://docs.cloud.google.com/gemini/docs/discover/data-governance
275. https://support.google.com/gemini/thread/327172086/can-you-provide-detailed-information-on-how-data-is-processed-and-stored-when-using-gemini-advanced?hl=en
276. https://note.com/ai_hikaku_/n/nc3ac13b4d7c4?hl=en
277. https://note.com/purin_games/n/n2269e806728e?hl=en
278. https://discuss.ai.google.dev/t/is-my-data-used-for-training-or-retained-with-gemini-paid-api-calls/64837
279. https://code.claude.com/docs/en/agent-sdk/structured-outputs
280. https://stackoverflow.com/questions/79543450/user-data-prompt-security-in-gcp
281. https://meetily.ai/llm-privacy/gemini
282. https://artificialanalysis.ai/models/gemini-3-8-flash
283. https://openrouter.ai/google/gemini-3.8-flash-20260902
284. https://artificialanalysis.ai/models/gemini-3-8-flash/providers
285. https://artificialanalysis.ai/models/gemini-3-8-flash-low/providers
286. https://codersera.com/blog/gemini-3-8-flash-complete-guide-2026/
287. https://intuitionlabs.ai/pdfs/gemini-3-8-flash-research-tasks-review.pdf
288. https://intuitionlabs.ai/articles/gemini-3-8-flash-research-tasks-review
289. https://x.com/ArtificialAnlys/status/2095177419489739063
290. https://docs.aws.amazon.com/claude-platform/latest/userguide/rate-limits.html
291. https://ground.news/article/google-gemini-38-flash-and-38-flash-cyber-presented
292. https://kilo.ai/models/openai-gpt-6-sol
293. https://benchlm.ai/models/gpt-6-sol
294. https://www.datacamp.com/blog/gpt-6-1-sol
295. https://emergent.sh/learn/gpt-6-sol-benchmarks
296. https://windowsforum.com/news/claude-sonnet-5-5-vs-gpt-6-sol-api-pricing-long-context-costs-and-coding-benchmarks.446361/
297. https://coursiv.io/blog/gpt-6-sol-luna
298. https://benchlm.ai/models/claude-haiku-5-5
299. https://www.explainx.ai/blog/claude-haiku-5-5-launch-pricing-benchmarks-sonnet-cache-cut-2026
300. https://magai.co/claude-haiku-5-5/
301. https://developers.openai.com/api/docs/guides/latest-model?model=gpt-5.2
302. https://kingy.ai/blog/gpt-sol-luna-vs-claude-haiku-5-5/
303. https://ppc.land/anthropic-haiku-5-5-costs-90-less-than-4-5-for-prompts-to-100-000-tokens/
304. https://aitoolsreview.co.uk/insights/gpt-6-sol-luna
305. https://docsbot.ai/models/compare/gpt-5-6-luna/gpt-6.1-sol
306. https://kingy.ai/blog/gpt-6-sol-luna-specs-benchmarks-pricing-comparison/
307. https://llm-stats.com/models/compare/gpt-6-luna-vs-gpt-6.1-sol
308. https://benchlm.ai/compare/gpt-6-1-sol-vs-gpt-6-luna-medium
309. https://www.vellum.ai/blog/gpt-6-1-sol-benchmarks-explained
310. https://github.com/joonlab/gpt-6.1-sol-benchmark
311. https://note.com/made_from_fuji/n/nc5d3a53e4434?hl=en
312. https://www.anthropic.com/news
313. https://docs.typesafe.ai/introduction
314. https://jev.pro/access/pricing/
315. https://jevtypesafe.org/privacy/
316. https://jevwiki.ai/wiki/reference/legal-and-data.md
317. https://typesafe-jev.com/en/faq/
318. https://jevwiki.ai/raw/wiki/reference/legal-and-data.md
319. https://gist.github.com/pjburnhill/adf8d28efcad9df037bfdece178ef965
320. https://docs.litellm.ai/blog/typesafe_jev
321. https://flaviocopes.com/jev-pricing/
322. https://www.developersdigest.tech/blog/typesafe-jev-system-one-models-release-guide-2026
323. https://openai.com/index/introducing-gpt-6-sol-and-luna/
324. https://huggingface.co/blog/sora-2/jev-ai-typesafe-ai-a-practical-guide-to-typed-deci
325. https://docs.perplexity.ai/docs/getting-started/pricing
326. https://documentation.decisions.com/docs/ai-openai

327. https://discuss.ai.google.dev/t/google-ai-studio-build-repeatedly-returns-quota-exceeded-please-try-again-later/185508
328. https://kie.ai/blog/what-is-gemini-3-8-flash
329. https://discuss.ai.google.dev/t/regarding-gemini-models-rate-limit-in-paid-tier/143557
330. https://bazaarlink.ai/en/blog/gemini-api-pricing-twd-2026-en
331. https://ai-api-hub.com/providers/google/rate-limits/
332. https://blog.laozhang.ai/en/posts/gemini-api-pricing
333. https://felloai.com/gemini-pricing/
334. https://developers.openai.com/api/docs/guides/decisions
335. https://platform.claude.com/docs/en/build-with-claude/structured-outputs
336. https://www.layer3labs.io/guides/gemini-3-8-flash-limits
337. https://book.st-hakky.com/en/data-science/gemini-pro-api-guide
338. https://klymentiev.com/blog/free-llm-api
339. https://aihubmix.com/blog/claude-haiku-5-5-vs-haiku-4-5-what-ten-cents-now-buys
340. https://llm-stats.com/blog/research/claude-haiku-5-5-launch
341. https://innfactory.ai/en/ai-models/anthropic-claude/
342. https://aicatchup.com/news/claude-haiku-5-5
343. https://www.techtimes.com/articles/328793/20261009/claude-haiku-55-becomes-viable-sub-agent-human-level-computer-use-ninety-percent-lower-cost.htm
344. https://help.openai.com/en/articles/6825453-chatgpt-release-notes
345. https://openai.com/news/product-releases/
346. https://developers.openai.com/api/docs/guides/latest-model?model=gpt-5.3-codex
347. https://go.tabbit.ai/model/gpt-6-sol/reviews/openai-official-gpt-6-sol-benchmarks
348. https://cellcog.ai/blog/gpt-6-sol-release-date/
349. https://emergent.sh/learn/gpt-6-luna-benchmarks
350. https://www.datacamp.com/blog/gpt-6-sol-and-luna
351. https://openai.github.io/openai-agents-js/guides/models/
352. https://openai.github.io/openai-agents-python/models/
353. https://community.openai.com/t/bug-gpt-5-6-luna-returns-500-with-presence-frequency-penalties-reasoning-effort-none/1395004
354. https://pypi.org/project/pydantic/
355. https://github.com/pydantic/pydantic-ai/actions/runs/37092926449/attempts/1
356. https://github.com/pydantic
357. https://openrouter.ai/anthropic/claude-haiku-5.5
358. https://github.com/pydantic/pydantic-core/releases
359. https://github.com/pydantic/pydantic-ai/blob/main/docs/changelog.md
360. https://tosea.ai/blog/gpt-6-sol-luna-complete-guide
361. https://openrouter.ai/anthropic/claude-sonnet-5.5
362. https://artificialanalysis.ai/models/creators/anthropic
363. https://artificialanalysis.ai/providers/anthropic
364. https://help.openai.com/en/articles/20001106-codex-rate-card
365. https://community.openai.com/t/announcing-gpt-6-sol-and-gpt-6-luna-in-the-api-codex-and-chatgpt/1399925
366. https://developers.openai.com/api/reference/python
367. https://pydantic.dev/docs/ai/api/pydantic-ai/capabilities/
368. https://openai.com/research/index/
369. https://github.com/pydantic/pydantic-ai/blob/main/docs/navigation.yml
370. https://pydantic.dev/pydantic-ai
371. https://pydantic.dev/
372. https://docs.venice.ai/es/guides/integrations/pydanticai
373. https://platform.claude.com/docs/en/build-with-claude/prompt-caching
374. https://apidog.com/blog/how-to-use-claude-haiku-5-5-api/
375. https://openrouter.ai/anthropic/claude-haiku-5.5/llms.txt
376. https://docs.openclaw.ai/providers/anthropic
377. https://www.cometapi.com/how-to-use-claude-haiku-5-5-api/
378. https://www.activepieces.com/blog/claude-haiku-55-api-pricing-rates-and-costs-for-2026
379. https://platform.claude.com/docs/en/build-with-claude/working-with-messages
380. https://ai.google.dev/gemini-api/docs/generate-content/thinking
381. https://googleapis.github.io/js-genai/release_docs/interfaces/types.UsageMetadata.html

382. https://docs.cloud.google.com/gemini-enterprise-agent-platform/reference/rest/v1/GenerateContentResponse
383. https://docs.cloud.google.com/gemini-enterprise-agent-platform/reference/models/inference
384. https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/capabilities/list-token
385. https://docs.nvidia.com/aiperf/dev/reference/vendor-usage-field-reference
386. https://docs.cloud.google.com/gemini-enterprise-agent-platform/models/capabilities/add-labels-to-api-calls
387. https://gemini.it-docs.cn/api/tokens.html?hl=zh-cn
388. https://github.com/google-gemini/deprecated-generative-ai-python/blob/main/docs/api/google/generativeai/protos/GenerateContentResponse/UsageMetadata.md
389. https://docs.aws.amazon.com/bedrock/latest/userguide/model-card-openai-gpt-6-1-sol.html
390. https://platform.openai.com/docs/changelog
391. https://community.openai.com/t/ultrafast-is-rolling-out-today-for-gpt-6-1-sol-in-the-api-codex-and-chatgpt-work/1404475
392. https://aws.amazon.com/blogs/machine-learning/bring-near-astra-intelligence-to-everyday-work-with-gpt-6-1-sol-on-amazon-bedrock/
393. https://help.openai.com/en/articles/8983778-chat-and-file-retention-policies-in-chatgpt
394. https://openai.com/index/devday-2026-recap/
395. https://help.kagi.com/kagi/ai/llms-privacy.html
396. https://www.layer3labs.io/guides/gpt-6-1-sol-review
397. https://www.layer3labs.io/guides/gpt-6-1-sol-benchmarks
398. https://emergent.sh/news/openai-devday-2026
399. https://note.com/duke_technology/n/nc9d39aaf2d7f?hl=en
400. https://aitoolsreview.co.uk/insights/gpt-6-1-sol
401. https://docs.anthropic.com/en/docs/about-claude/models/overview
402. https://www.datacamp.com/blog/gpt-6-1-sol-vs-opus-5-5
403. https://pydantic.dev/docs/ai/integrations/logfire/
404. https://pydantic.dev/articles/vercel-ai-sdk-logfire-otel
405. https://docs.base14.io/instrument/apps/auto-instrumentation/pydantic-ai/
406. https://pypi.org/project/genai-otel-instrument/
407. https://deepwiki.com/pydantic/pydantic-ai/5-observability-and-instrumentation
408. https://logfire.pydantic.dev/docs/guides/web-ui/llm-panels/
409. https://deepwiki.com/pydantic/pydantic-ai/5.1-instrumentation-architecture
410. https://github.com/Kludex/pydantic-ai-go/blob/main/docs/observability.md
411. https://developers.openai.com/api/docs/guides/tools
412. https://docs.aws.amazon.com/ja_jp/claude-platform/latest/userguide/rate-limits.md
413. https://help.openai.com/articles/10478918
414. https://help.openai.com/en/articles/6614209-how-do-i-check-my-token-usage
415. https://help.openai.com/en/articles/4936856-what-are-tokens-and-how-to-count-them
416. https://vercel.com/docs/ai-gateway/modalities/decision
417. https://github.com/stacklok/mecatl/blob/main/docs/adr/0017-openai-responses-api.md
418. https://docs.litellm.ai/docs/decisions
419. https://platform.claude.com/docs/en/api/overview
420. https://platform.claude.com/docs/en/build-with-claude/handling-stop-reasons
421. https://platform.claude.com/docs/en/api/messages
422. https://platform.claude.com/docs/en/api/beta/messages
423. https://developers.openai.com/api/docs/changelog.md
424. https://platform.claude.com/docs/en/release-notes/overview
425. https://platform.claude.com/docs/en/build-with-claude/token-counting.md
426. https://platform.claude.com/docs/en/api/messages/create
427. https://platform.claude.com/docs/en/api/messages/count_tokens
428. https://platform.claude.com/docs/en/api/cli/beta/messages/count_tokens
429. https://platform.claude.com/docs/en/api/csharp/messages/count_tokens
430. https://langwatch.ai/models/gemini-3-8-flash
431. https://benchmarks.speko.ai/llm/gemini-gemini-3-8-flash
432. https://www.callmissed.com/blog/claude-opus-vs-gemini-flash-api-buyer
433. https://llm-stats.com/models/gemini-3.8-flash
434. https://pydantic.dev/docs/ai/models/anthropic/
435. https://kingy.ai/blog/gemini-3-8-flash-review-benchmarks-specs-pricing/
436. https://local-ai-zone.github.io/blog/Gemini_3.8_Flash_Technical_Breakdown.html

437. https://artificialanalysis.ai/models/releases/gpt-6-luna
438. https://alphasignal.ai/news/openai-s-decisions-api-promises-routing-answers-10x-faster-for-developers
439. https://artificialanalysis.ai/models/claude-haiku-5-5-low
440. https://artificialanalysis.ai/models/claude-haiku-5-5-low/providers
441. https://www.moneycontrol.com/technology/anthropic-launches-claude-haiku-5-5-its-most-affordable-and-fastest-model-yet-article-14047297.html
442. https://timesofindia.indiatimes.com/technology/tech-news/anthropic-launches-claude-haiku-5-5-its-fastest-and-cheapest-ai-model-yet/articleshow/134780000.cms
443. https://artificialanalysis.ai/models/gpt-6-luna
444. https://artificialanalysis.ai/models/gpt-6-luna-non-reasoning/providers
445. https://docs.langchain.com/langsmith/llm-gateway-decision-models
446. https://platform.claude.com/docs/en/api/service-tiers
447. https://artificialanalysis.ai/providers/openai
448. https://artificialanalysis.ai/models/gpt-6-luna-low/providers
449. https://artificialanalysis.ai/models/gpt-6-luna/providers
450. https://artificialanalysis.ai/models/gpt-6-luna-high/providers
451. https://blog.google/innovation-and-ai/models-and-research/gemini-models/3-8-flash-and-3-8-flash-cyber/
452. https://developers.openai.com/cookbook/examples/responses_api/reasoning_items
453. https://developers.openai.com/api/docs/guides/migrate-to-responses
454. https://help.openai.com/articles/4936856
455. https://github.com/openai/openai-python/blob/main/src/openai/types/responses/response_usage.py
456. https://community.openai.com/t/usage-info-in-api-responses/18862
457. https://pydantic.dev/docs/ai/overview/
458. https://developers.openai.com/api/docs/models/gpt-5.6-luna
459. https://developers.openai.com/cookbook/examples/how_to_handle_rate_limits
460. https://help.openai.com/en/articles/6643435-how-do-i-get-more-tokens-or-increase-my-monthly-usage-limits
461. https://help.openai.com/en/articles/5955604-troubleshooting-api-rate-limits-and-429-errors
462. https://help.openai.com/en/articles/6891753-what-are-the-best-practices-for-managing-my-rate-limits-in-the-api%25253F.doc
463. https://www.aifreeapi.com/en/posts/openai-decisions-api
464. https://note.com/0x__tom/n/n85c5716ce442?hl=en
465. https://x.com/i/trending/2107091883243827397
466. https://www.eesel.ai/blog/openai-decisions-api
467. https://note.com/prodready_ai/n/n3b4d2d416bdf?hl=en
468. https://platform.claude.com/docs/en/api/rate-limits
469. https://platform.claude.com/docs/en/resources/overview
470. https://note.com/s_syn25/n/n15894dbc45e2?hl=en
471. https://pydantic.dev/docs/ai/api/models/base/
472. https://developers.openai.com/api/docs/models/gpt-6-sol
473. https://ai.google.dev/gemini-api/docs/pricing
474. https://ai.google.dev/gemini-api/docs/models/gemini-3-flash-preview
475. https://developers.openai.com/api/docs/models/gpt-6.1-sol
476. https://pydantic.dev/docs/ai/models/openai/
477. https://docs.anthropic.com/en/docs/about-claude/models/whats-new-claude-4-8
478. https://ai.google.dev/gemini-api/docs/models/gemini-3.7-flash
479. https://ai.google.dev/gemini-api/docs/models
480. https://developers.openai.com/api/docs/pricing
481. https://cdn.openai.com/pdf/gpt-6-october.pdf
482. https://ai.google.dev/gemini-api/docs/rate-limits
483. https://aistudio.google.com/models/gemini-3
484. https://developers.openai.com/api/docs/models
485. https://openai.com/safety/evaluations-hub/
486. https://www.anthropic.com/document/claude-haiku-5-5-system-card
487. https://docs.anthropic.com/en/release-notes/api
488. https://ai.google.dev/gemini-api/docs/function-calling
489. https://ai.google.dev/gemini-api/docs/gemini-3
490. https://ai.google.dev/gemini-api/docs/latest-model

491. https://ai.google.dev/gemini-api/docs/models/gemini-3.8-flash
492. https://ai.google.dev/gemini-api/docs/zdr
493. https://developers.openai.com/api/docs/models/gpt-6-sol.md
494. https://ai.google.dev/gemini-api/terms-archive/terms_05_02_24
495. https://ai.google.dev/gemini-api/docs/generate-content/latest-model
496. https://ai.google.dev/gemini-api/terms-archive/terms_12_20_24
497. https://ai.google.dev/gemini-api/docs/changelog
498. https://developers.cloudflare.com/ai/models/openai/gpt-6-luna/
499. https://openai.com/index/introducing-gpt-6-1-sol/
500. https://ai.google.dev/gemini-api/terms.md.txt
501. https://deepmind.google/models/gemini/flash/
502. https://ai.google.dev/gemini-api/docs/models/gemini-3.5-flash
503. https://deepmind.google/models/model-cards/gemini-3-8-flash/
504. https://ai.google.dev/gemini-api/docs/logs-policy
505. https://artificialanalysis.ai/models/releases/gemini-3-8-flash
506. https://developers.cloudflare.com/ai/models/typesafe/jev/index.md
507. https://ai.google.dev/gemini-api/docs/rate-limits?hl=en
508. https://pydantic.dev/docs/ai/models/typesafe/
509. https://developers.openai.com/api/docs/guides/decisions-voice
510. https://developers.openai.com/api/docs/guides/your-data.md
511. https://ai.google.dev/gemini-api/docs/pricing.md.txt
512. https://apidog.com/blog/claude-haiku-5-5-benchmarks/
513. https://itbrief.com.au/story/openai-launches-gpt-6-1-sol-as-cheaper-coding-model
514. https://docs.typesafe.ai/models
515. https://geminicli.com/docs/resources/quota-and-pricing/
516. https://heymark.ai/en/guides/claude-haiku-5-5-benchmarks-pricing/
517. https://computingforgeeks.com/gpt-6-sol-luna-released-features-benchmarks/
518. https://typesafe.ai/enterprise
519. https://apidog.com/blog/gemini-3-8-flash-pricing/